

Enable Intel Control Flow Enforcement Technology

Chen Hu <hu1.chen@intel.com>

Intel System Software Engineering

Oct. 2022



Agenda


- Introduction
 - Shadow Stack (SS)
 - Indirect Branch Tracking (IBT)
- SW stack to enable CET
 - App Execution Flow
 - Kernel Implementation and challenges
- User Space CET Reference Design
- CET Perf Impact
- Q & A

Introduction – Control Flow

Return/jump oriented programming (ROP) attack here.

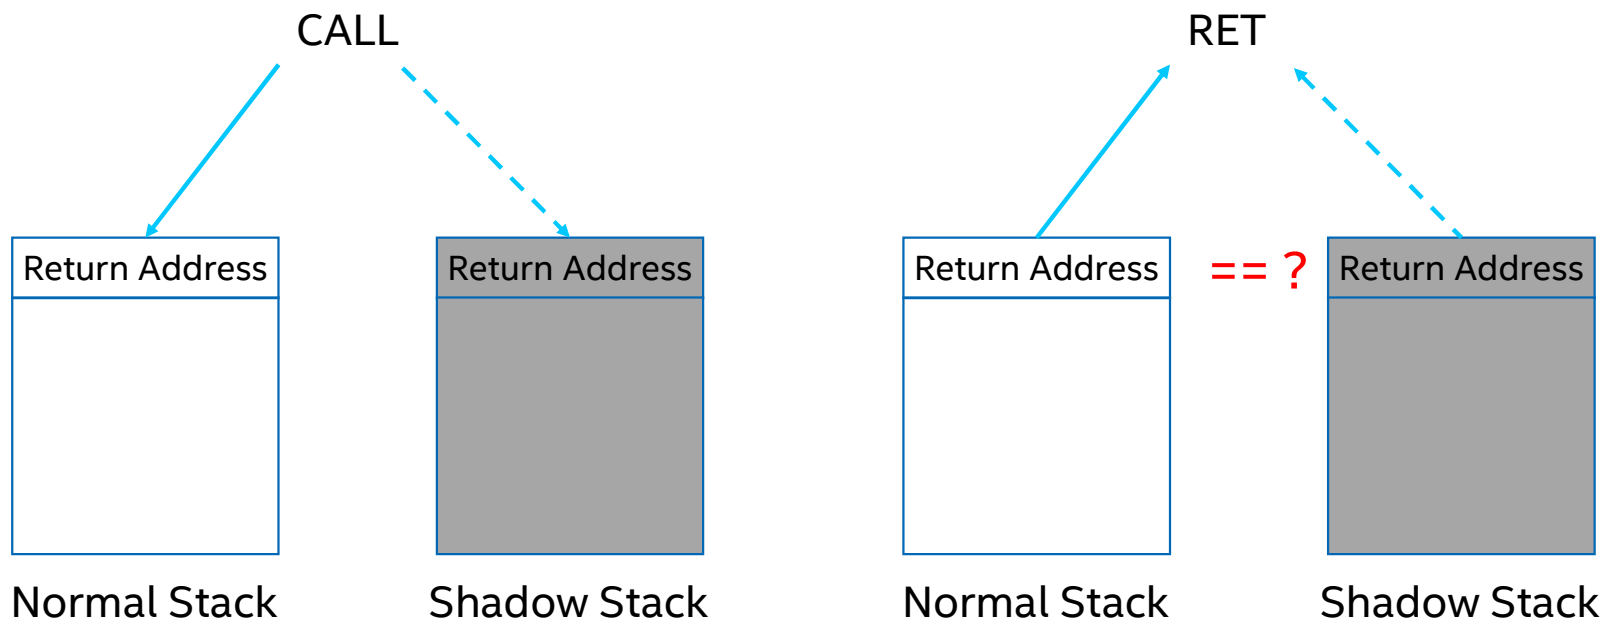


Intel Control Flow Enforcement (CET) mitigates this.



Control Flow	Mitigation
CALL / RET	Shadow Stack
JMP *%rax	Indirect Branch Tracking (IBT)

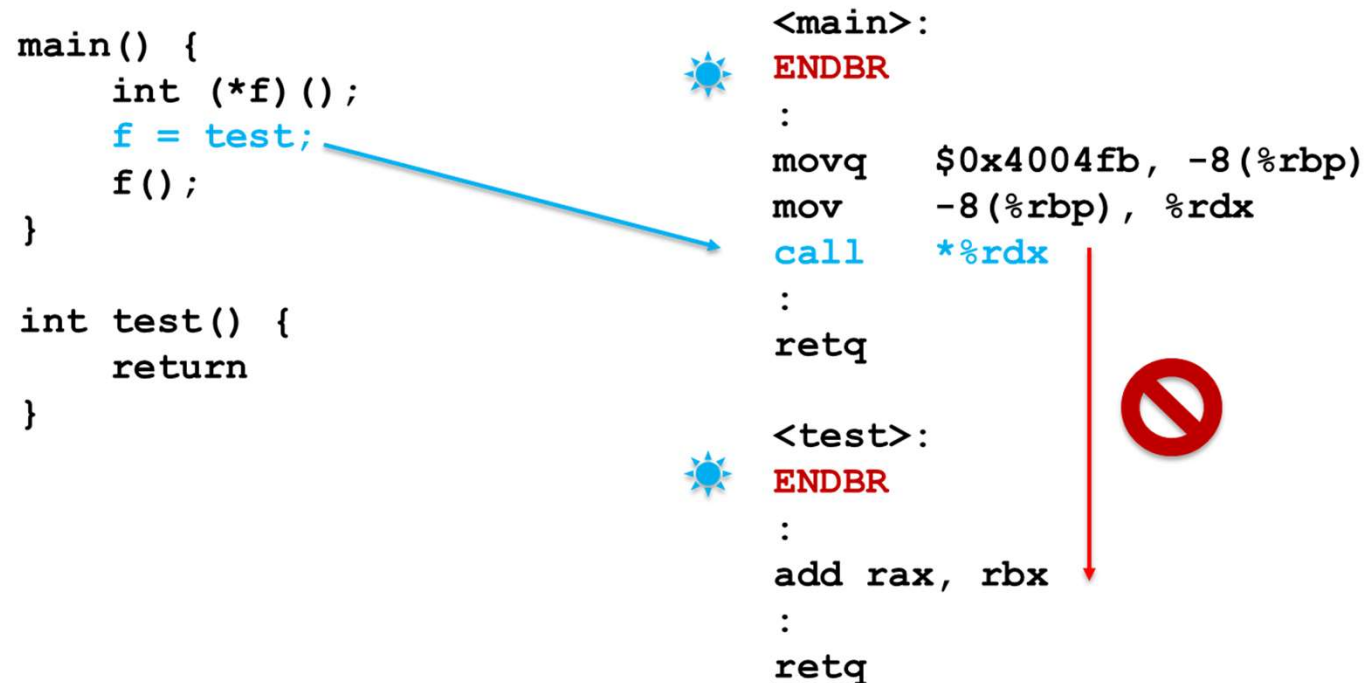
CET – Shadow Stack



- Return addresses popped off both stacks should match or suffer #CP

Credit: hongjiu.lu@intel.com <https://lpc.events/event/2/contributions/147/attachments/72/83/CET-LPC-2018.pdf>

CET – Indirect Branch Tracking



- Any indirect CALL/JMP must target an ENDBR instruction or suffer #CP
- ENDBR is NOP on non-CET processor

SW Stack to Enable CET

Compile		Run	
Compiler		Application	
Linker		Libc	
		Kernel	IBT for kernel
			IBT for user
			SS for kernel
			SS for user

More complex than imagined!

CET violation example

- Compile with “-fcf-protection=full/branch/ret/none”
- ELF header indicates if CET enabled

```
sdp@b49691a74be4:~/chen/php-src$ sudo readelf -n ./sapi/cli/php

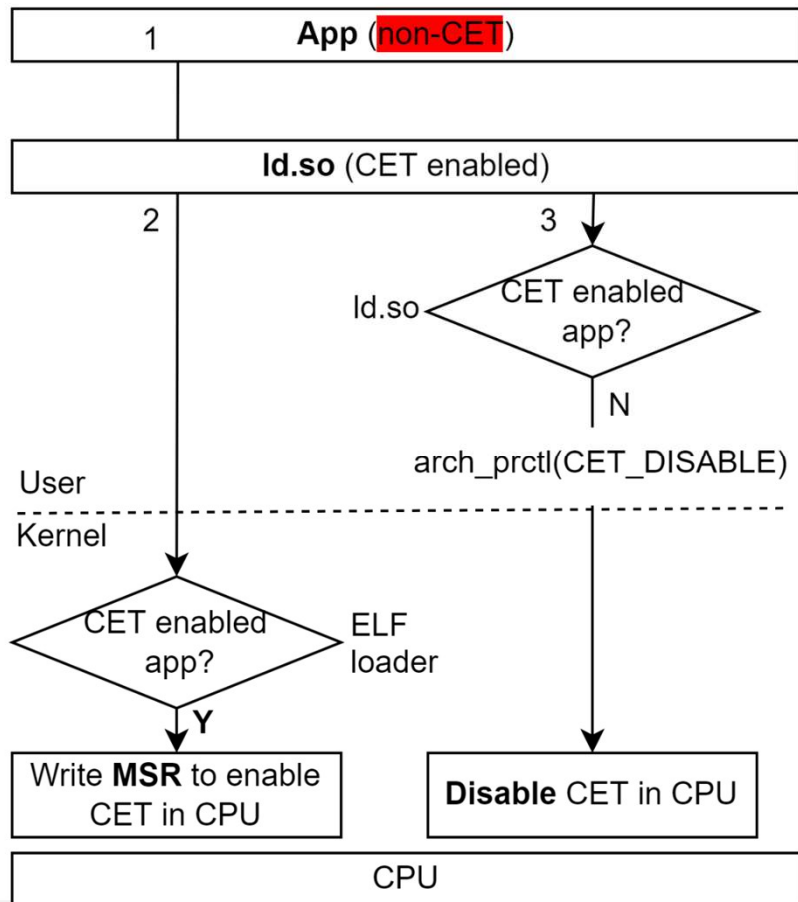
Displaying notes found in: .note.gnu.property
  Owner                Data size  Description
  GNU                  0x00000020  NT_GNU_PROPERTY_TYPE_0
    Properties: x86 feature: IBT, SHSTK
    x86 ISA needed: x86-64-baseline
```

- Suffer #CP while CET is violated

```
sdp@b49691a74be4:~/chen/php-src$ ./sapi/cli/php -d zend_extension=/home/sdp/chen/php-src/modules/opcache.so -d opcache.enable=1 -d opcache.enable_cli=1 -d opcache.jit_buffer_size=128M -d opcache.jit=tracing test.php
Segmentation fault
sdp@b49691a74be4:~/chen/php-src$
sdp@b49691a74be4:~/chen/php-src$ dmesg
[3400105.006718] traps: php[2871286] control protection ip:49f68670 sp:7ffc44c98840 ssp:7f740b7f3fc0 error:3(endbranch) in zero (deleted)[49f68000+8000000]
```

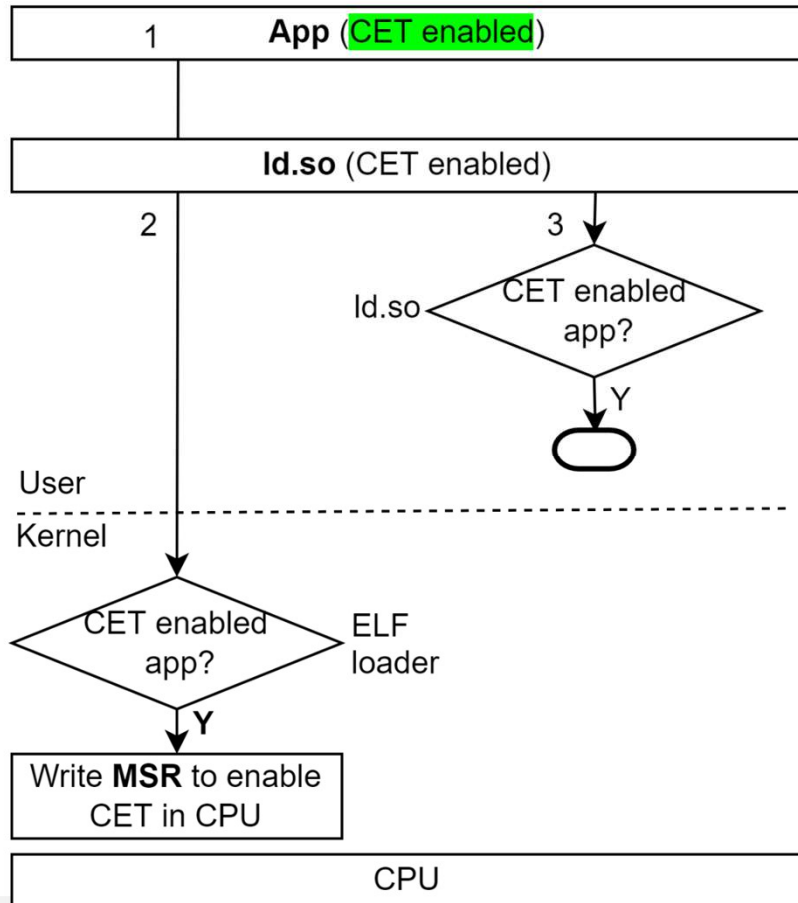
Ubuntu22.04 on Intel Sapphire Rapid with 5.15 kernel

Workflow of non-CET App



1. ELF header of app indicates non-CET.
2. Since ld.so is CET enabled, all apps start as CET enabled. Kernel set MSR to enable CET.
3. ld.so find that app is non-CET. Issue syscall to disable CET.

Workflow of CET Enabled App



1. ELF header of app indicates CET enabled.
2. Since ld.so is CET enabled, all apps start as CET enabled. Kernel set MSR to enable CET.
3. ld.so find that app is CET enabled. Do nothing.

Kernel Implementation – Shadow Stack

\$./foo

load_elf_binary

arch_setup_elf_property

shstk_setup:

1. Allocate shadow stack
2. Set SS pointer:


MSR IA32_PL3_SSP

3. Enable shadow stack:

MSR IA32_U_CET

- Create process

Kernel Mem



- Context Switch

Kernel Implementation – Challenges

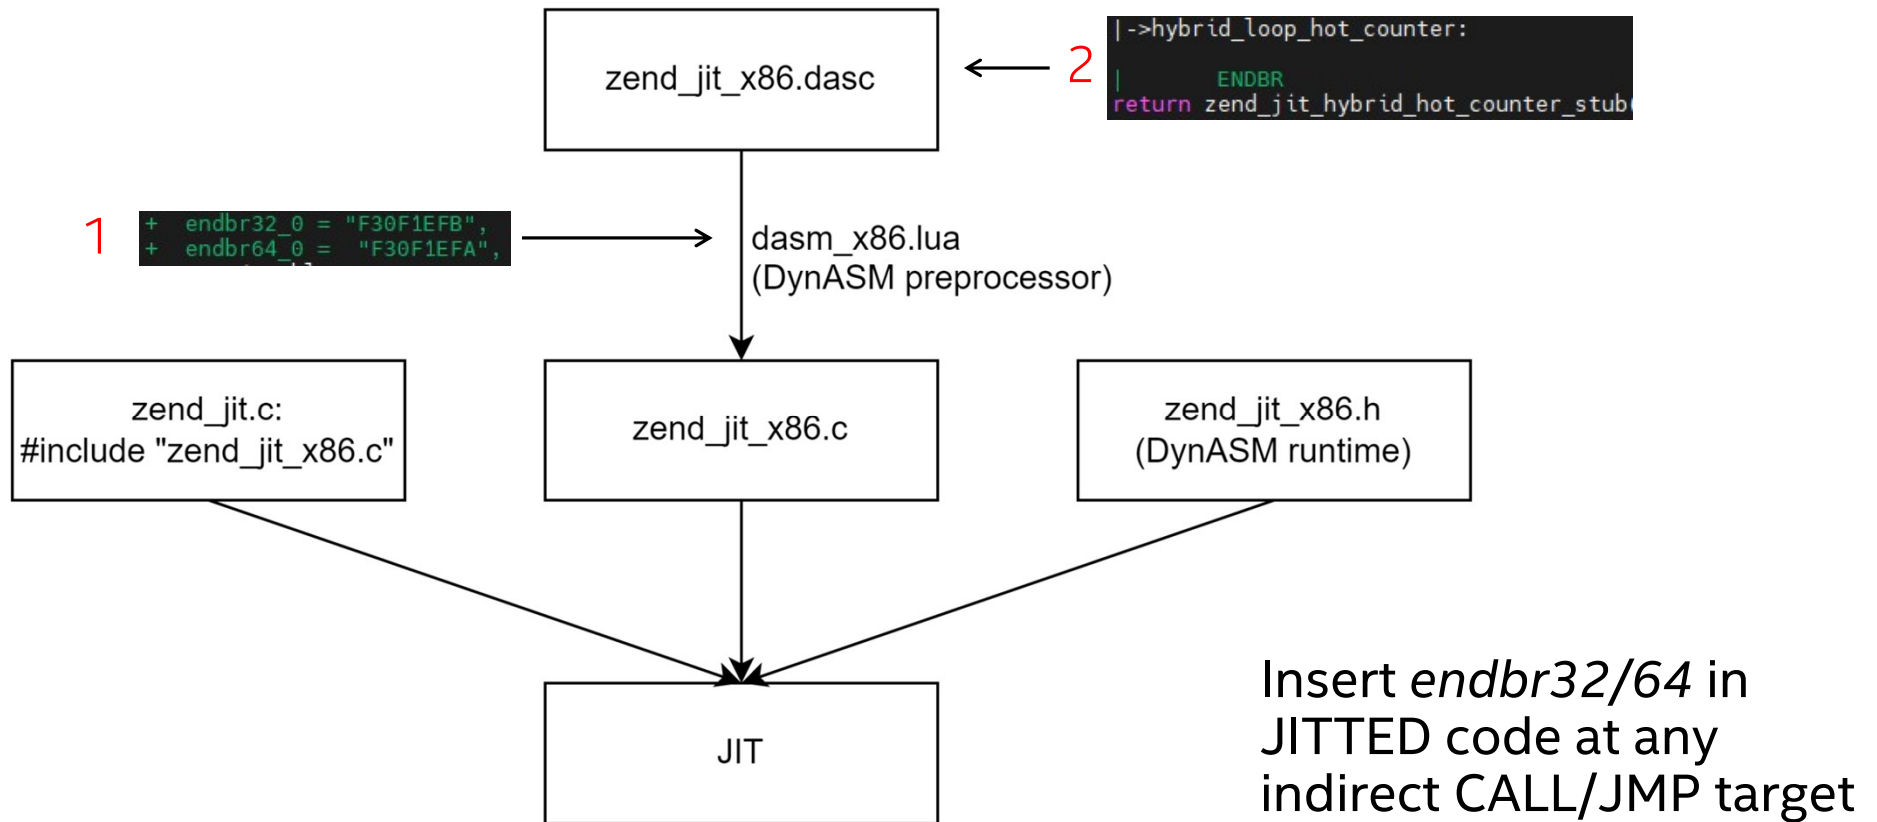
1. Shadow stack has “write=0, dirty=1 PTEs”
2. Signal
3. ABI issues
4. CRIU (checkpoint and restore in user space) support

Status

Compiler	GCC 8	Application	See Ref
Linker	Binutils 2.31	Libc	Glibc 2.28
		Kernel	IBT for kernel
			IBT for user
			SS for kernel
			SS for user ¹

1. <https://lkml.org/lkml/2022/9/29/1149>

User Space CET Ref Design: PHP JIT



User Space CET Ref Design: PHP JIT

Php:

```
Dump of assembler code from 0x48000630 to 0x48000660:  
0x0000000048000630 <JIT$$hybrid_hot_trace+0>: 66 c7 02 13 7f mov WORD PTR [rdx],0x7f13  
0x0000000048000635 <JIT$$hybrid_hot_trace+5>: 4c 89 f7 mov rdi,r14  
0x0000000048000638 <JIT$$hybrid_hot_trace+8>: 4c 89 fe mov rsi,r15  
0x000000004800063b <JIT$$hybrid_hot_trace+11>: 48 b8 60 4b 25 f5 ff 7f 00 00 movabs rax,0x7ffff5254b60
```




Php: Enable CET-IBT for JIT:

```
Dump of assembler code from 0x48000630 to 0x48000660:  
0x0000000048000630 <JIT$$hybrid_hot_trace+0>: f3 0f 1e fa endbr64  
0x0000000048000634 <JIT$$hybrid_hot_trace+4>: 66 c7 02 13 7f mov WORD PTR [rdx],0x7f13  
0x0000000048000639 <JIT$$hybrid_hot_trace+9>: 4c 89 f7 mov rdi,r14  
0x000000004800063c <JIT$$hybrid_hot_trace+12>: 4c 89 fe mov rsi,r15  
0x000000004800063f <JIT$$hybrid_hot_trace+15>: 48 b8 60 4b 25 f5 ff 7f 00 00 movabs rax,0x7ffff5254b60
```

<https://github.com/php/php-src/pull/8774>


User Space CET Ref Design: CoRoutine

map_shadow_stack()



1. Allocate shadow stack for new coroutine.

Normal Stack



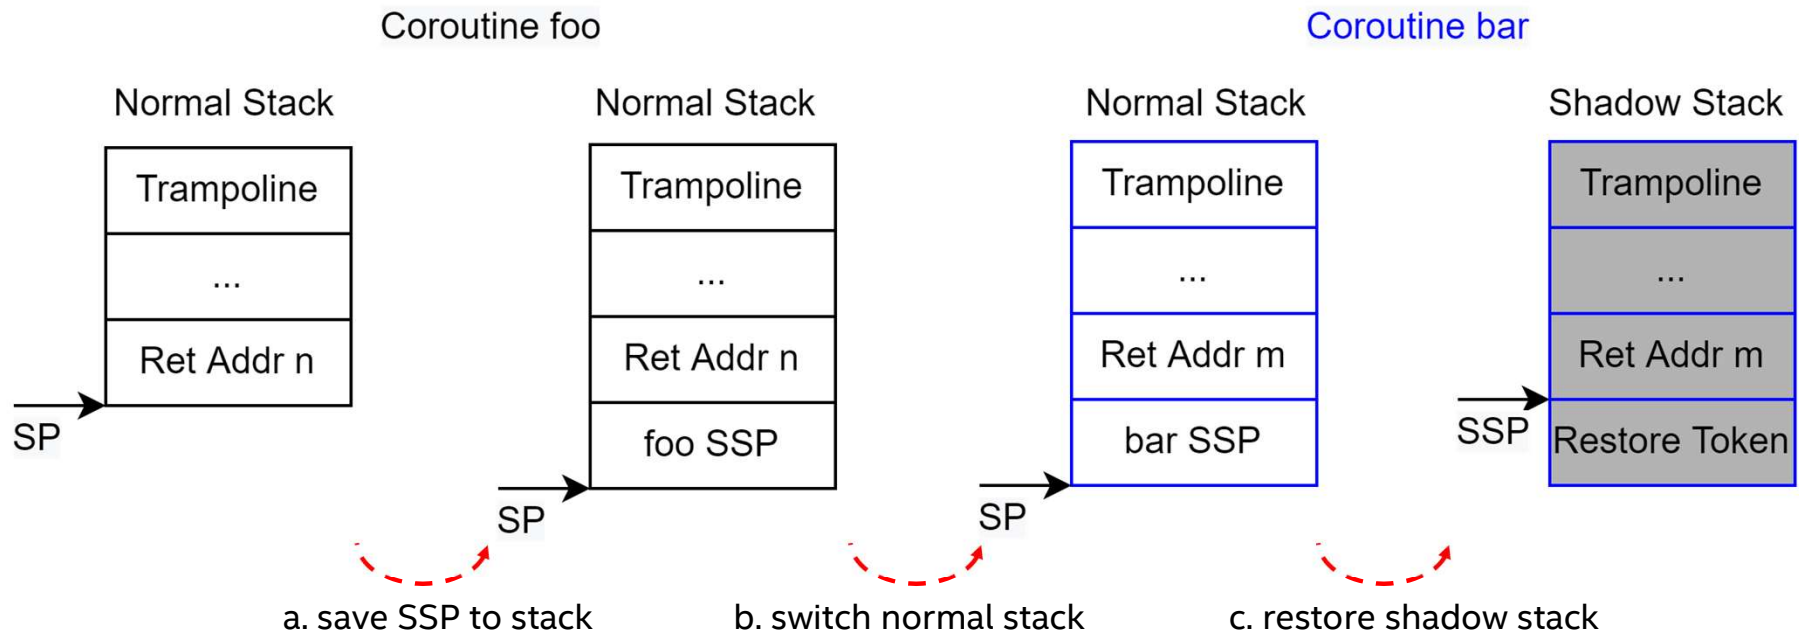
Shadow Stack



2. Populate normal / shadow stack when create fcontext

<https://github.com/boostorg/context/pull/207/>

User Space CET Ref Design: CoRoutine



3. Context switch from coroutine foo to bar

CET Perf Impact

- IBT
 - More uops used
 - More instructions (endbr) in binary.
 - If CPU doesn't support IBT or doesn't enable IBT, endbr is nop.
- Shadow stack
 - More uops used
 - Two additional mem access (push/pop shadow stack) for each function call
 - Compare normal / shadow stack during return
- Speculation behavior (Refer to Intel SDM)

Acknowledgement

Srinivas Suresh, H.J. Lu, Rick Edgecombe, Yu-cheng Yu, Yuhan Yang

Q & A

